

Spatial Analysis of Madrid Traffic Data with MetricGraph

David Bolin and Alexandre Simas - INLA Workshop 2025

May 2025

1 Introduction

This tutorial explores spatial analysis of traffic data on road networks using the **MetricGraph** package. We'll work with real traffic data from Madrid, Spain, and learn how to:

1. Construct a metric graph from OpenStreetMap road network data
2. Add observations to the graph
3. Visualize traffic data on the network
4. Fit spatial models to analyze traffic patterns

The **MetricGraph** package provides tools for statistical analysis of data on network-structured domains, like road networks. It implements various types of Gaussian random fields on graphs and offers interfaces to INLA for Bayesian inference.

Let's begin by loading the necessary packages:

```
library(data.table)
library(sf)
library(osmdata)
library(dplyr)
library(MetricGraph)
library(ggplot2)
library(inlabru)
```

2 Data Description

For this tutorial, we've prepared datasets containing traffic information for Madrid. The datasets include:

1. Traffic sensor locations in Madrid
2. Traffic measurements from these sensors during evening rush hour (6-7 PM) in March 2025

These data have been generated based on typical traffic patterns and distributions for educational purposes.

2.1 Loading the Data

For convenience, we've prepared data files containing: - The traffic sensor locations - The traffic measurements

Now we can load the pre-prepared data:

```
``` r
Load the pre-prepared data files
madrid_traffic_avg_sf <- readRDS("madrid_traffic_avg.rds")
radar_locations_sf <- readRDS("radar_locations.rds")
```

```
Examine the data
head(madrid_traffic_avg_sf)

Simple feature collection with 6 features and 4 fields
Geometry type: POINT
Dimension: XY
Bounding box: xmin: -3.666633 ymin: 40.40367 xmax: -3.663641 ymax: 40.40921
Geodetic CRS: WGS 84
A tibble: 6 x 5
geometry avg_speed intensity occupation load
<POINT [°]> <dbl> <dbl> <dbl> <dbl>
1 (-3.664111 40.40921) 73 3064. 15 0
2 (-3.663641 40.40657) 58.4 3160. 17.5 0
3 (-3.666633 40.40434) 53 1075. 12.3 0
4 (-3.665753 40.40372) 50.0 1118. 10.2 0
5 (-3.666589 40.40367) 68.8 1395. 5.6 0
6 (-3.665522 40.40554) 62.4 1476. 6.2 0
```

## 3 Madrid Road Network as a Metric Graph

### 3.1 Understanding the Study Area

For this tutorial, we're focusing on a central area of Madrid defined by this bounding box:

```
Define the Madrid bounding box used for this tutorial
small_bbox <- c(-3.73, 40.40, -3.65, 40.45)

Create a simple data frame for the bounding box to visualize it
bbox_df <- data.frame(
 lon = c(small_bbox[1], small_bbox[3], small_bbox[3], small_bbox[1], small_bbox[1]),
 lat = c(small_bbox[2], small_bbox[2], small_bbox[4], small_bbox[4], small_bbox[2])
)
bbox_sf <- st_as_sf(bbox_df, coords = c("lon", "lat"), crs = 4326) %>%
 st_combine() %>%
 st_cast("POLYGON")
```

### 3.2 Exercise 1: Exploring the Graph Structure

Let us start by creating a graph from the OpenStreetMap data with some standard features.

```
madrid_sf_filtered <- opq(bbox = small_bbox) %>%
 add_osm_feature(key = "highway",
 value = c("motorway", "trunk", "primary", "secondary")) %>%
 osmdata_sf()

madrid_graph_components <- graph_components$new(madrid_sf_filtered, perform_merges = TRUE,
 tolerance = list(vertex_vertex = 0.001,
 vertex_edge = 0.01))

madrid_graph <- madrid_graph_components$get_largest()
```

Let's explore the pre-loaded graph structure:

```
Look at the graph structure
print(madrid_graph)
```

```

A metric graph with 1328 vertices and 1496 edges.
##
Vertices:
Degree 1: 54; Degree 2: 981; Degree 3: 199; Degree 4: 91; Degree 5: 3;
With incompatible directions: 6
##
Edges:
Lengths:
Min: 0.001153276 ; Max: 1.256895 ; Total: 145.4731
Weights:
Columns: osm_id name access access:lanes admin_level alt_name bicycle bicycle:backward:conditi
That are circles: 0
##
Graph units:
Vertices unit: degree ; Lengths unit: km
##
Longitude and Latitude coordinates: TRUE
Which spatial package: sp
CRS: +proj=longlat +datum=WGS84 +no_defs
##
Some characteristics of the graph:
Connected: TRUE
Has loops: FALSE
Has multiple edges: FALSE
Is a tree: FALSE
Distance consistent: unknown
Has Euclidean edges: unknown

What is the total length of the road network in the graph?
cat("Total length of road network:", round(sum(madrid_graph$edge_lengths), 2), "km\n")

Total length of road network: 145.47 km

Compute additional characteristics
madrid_graph$compute_characteristics()
print(t(madrid_graph$characteristics))

has_loops connected has_multiple_edges is_tree
[1,] FALSE TRUE FALSE FALSE

```


### 3.3 Visualizing the Graph

Let's visualize our road network graph:

```

Plot the graph
p <- madrid_graph$plot(vertex_size = 0.5, edge_width = 0.5)
print(p)

```



## 4 Traffic Data on the Road Network


### 4.1 Examining Sensor Locations

Let's examine the traffic sensor locations in our study area:

```
How many sensors are in our study area?
cat("Number of sensors in the study area:", nrow(radar_locations_sf), "\n")

Number of sensors in the study area: 1299

Plot the graph with the sensor locations
p <- madrid_graph$plot(vertex_size = 0, edge_width = 0.5)
p + geom_sf(data = radar_locations_sf, color = "red", size = 1)
```



## 4.2 Recreate the graph with more features

Let us recreate the graph with more features.

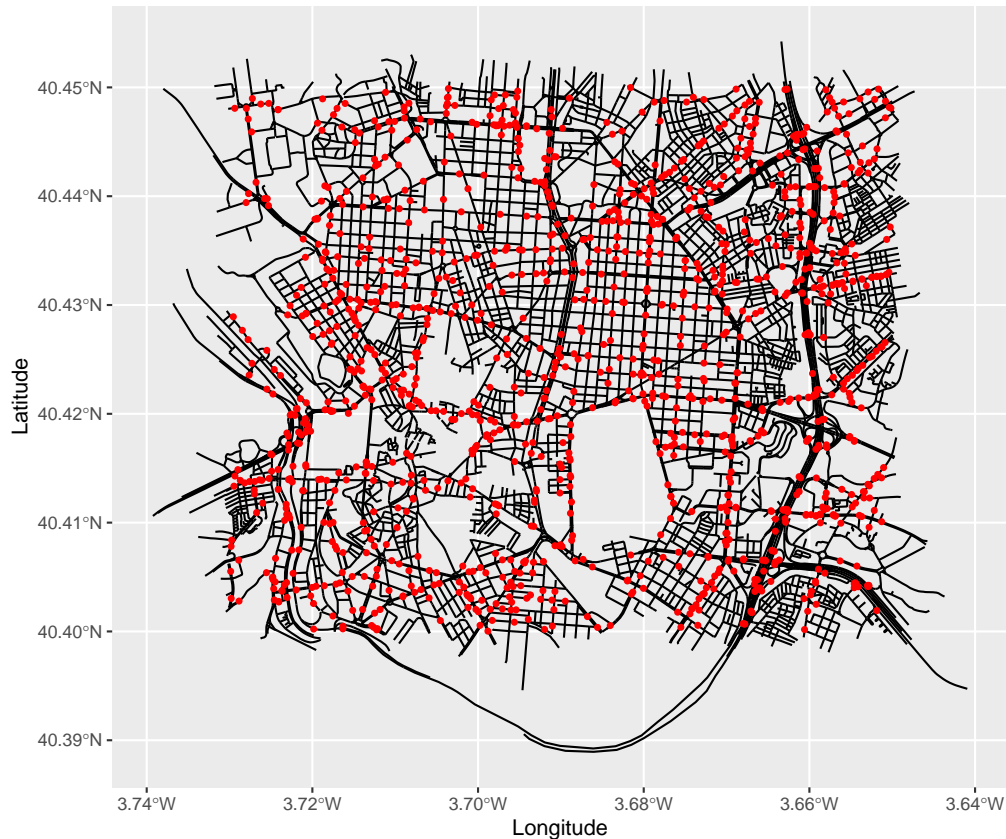
```
madrid_sf_new <- opq(bbox = small_bbox) %>%
 add_osm_feature(key = "highway",
 value = c("motorway", "trunk", "primary", "secondary",
 "tertiary", "residential", "unclassified",
 "motorway_link", "trunk_link")) %>%
 osmdata_sf()

madrid_graph_components_new <- graph_components$new(madrid_sf_new, perform_merges = TRUE, tolerance = 10)

madrid_graph_new <- madrid_graph_components_new$get_largest()
```

## 4.3 Examining the new graph along with the data

```
Plot the new graph with the sensor locations
p <- madrid_graph_new$plot(vertex_size = 0, edge_width = 0.5)
p + geom_sf(data = radar_locations_sf, color = "red", size = 1)
```



We can now add the data to the graph.

```
madrid_graph_new$add_observations(madrid_traffic_avg_sf)
```

```
$removed
[1] avg_speed intensity occupation load .coord_x .coord_y
<0 rows> (or 0-length row.names)
```

## 4.4 Exercise 2: Checking Coordinate Reference Systems

A crucial step in spatial analysis is ensuring that all data uses the same coordinate reference system (CRS):

```
Check the CRS of the radar locations
st_crs(radar_locations_sf)
```

```
Coordinate Reference System:
User input: EPSG:4326
wkt:
GEOGCRS["WGS 84",
ENSEMBLE["World Geodetic System 1984 ensemble",
MEMBER["World Geodetic System 1984 (Transit)"],
MEMBER["World Geodetic System 1984 (G730)"],
MEMBER["World Geodetic System 1984 (G873)"],
MEMBER["World Geodetic System 1984 (G1150)"],
MEMBER["World Geodetic System 1984 (G1674)"],
MEMBER["World Geodetic System 1984 (G1762)"],
MEMBER["World Geodetic System 1984 (G2139)"],
MEMBER["World Geodetic System 1984 (G2296)"],
ELLIPSOID["WGS 84",6378137,298.257223563,
```

```

LENGTHUNIT["metre",1]],
ENSEMBLEACCURACY[2.0]],
PRIMEM["Greenwich",0,
ANGLEUNIT["degree",0.0174532925199433]],
CS[ellipsoidal,2],
AXIS["geodetic latitude (Lat)",north,
ORDER[1],
ANGLEUNIT["degree",0.0174532925199433]],
AXIS["geodetic longitude (Lon)",east,
ORDER[2],
ANGLEUNIT["degree",0.0174532925199433]],
USAGE[
SCOPE["Horizontal component of 3D system."],
AREA["World."],
BBOX[-90,-180,90,180]],
ID["EPSG",4326]]

Check the CRS of the graph
print(madrid_graph)

A metric graph with 1328 vertices and 1496 edges.
##
Vertices:
Degree 1: 54; Degree 2: 981; Degree 3: 199; Degree 4: 91; Degree 5: 3;
With incompatible directions: 6
##
Edges:
Lengths:
Min: 0.001153276 ; Max: 1.256895 ; Total: 145.4731
Weights:
Columns: osm_id name access access:lanes admin_level alt_name bicycle bicycle:backward:conditi
That are circles: 0
##
Graph units:
Vertices unit: degree ; Lengths unit: km
##
Longitude and Latitude coordinates: TRUE
Which spatial package: sp
CRS: +proj=longlat +datum=WGS84 +no_defs
##
Some characteristics of the graph:
Connected: TRUE
Has loops: FALSE
Has multiple edges: FALSE
Is a tree: FALSE
Distance consistent: unknown
Has Euclidean edges: unknown

```

## 4.5 Examining Traffic Data

Now let's examine the pre-computed traffic measurements:


```

Basic statistics of the traffic data
summary(madrid_traffic_avg_sf)

```

```
geometry avg_speed intensity occupation load
POINT :1216 Min. : 0.000 Min. : 0.0 Min. : 0.00 Min. : 0.00
epsg:4326 : 0 1st Qu.: 0.000 1st Qu.: 258.2 1st Qu.: 4.90 1st Qu.: 20.35
+proj=long...: 0 Median : 0.000 Median : 515.8 Median : 8.10 Median : 32.90
Mean : 6.071 Mean : 788.8 Mean :11.83 Mean : 34.74
3rd Qu.: 0.000 3rd Qu.:1003.5 3rd Qu.:14.04 3rd Qu.: 45.35
Max. :90.800 Max. :6791.9 Max. :99.05 Max. :100.00
```

```
Histogram of traffic intensity
hist(madrid_traffic_avg_sf$intensity,
 main = "Distribution of Traffic Intensity",
 xlab = "Intensity (vehicles per 15min)",
 col = "lightblue")
```



## 5 Visualizing Traffic Data on the Graph

Let's check if the traffic data is already loaded into the graph:

```
madrid_graph_new$add_observations(madrid_traffic_avg_sf)


$removed
[1] avg_speed intensity occupation load .coord_x .coord_y
<0 rows> (or 0-length row.names)

madrid_graph_new$add_observations(
 madrid_graph_new$mutate(log_intensity = log(pmax(1, intensity))),
 clear_obs = TRUE
)

Plot the data
madrid_graph_new$plot(data = "log_intensity",
```



```
vertex_size = 0,
data_size = 2,
edge_width = 0.5) +
scale_color_viridis_c(name = "Log Intensity")
```




## 6 Exploratory Data Analysis

Let's explore the relationships between different traffic metrics:

```
Get the data from the graph
traffic_data <- madrid_graph_new$get_data()
Histogram of log intensity
hist(traffic_data$log_intensity,
 main = "Histogram of Log Traffic Intensity",
 xlab = "Log Intensity",
 col = "lightblue",
 breaks = 20)
```

**Histogram of Log Traffic Intensity**



## 7 Spatial Modeling

Now let's fit a spatial model to the log intensity using a Whittle-Matérn random field with  $\alpha = 1$ , which corresponds to an exponential covariance model:

```
Fit the model
fit_alpha1 <- graph_lme(log_intensity ~ 1,
 graph = madrid_graph_new,
 BC = 0,
 model = list(type = "WhittleMatern", alpha = 1))

Look at the model summary
summary(fit_alpha1)
```

```
##
Latent model - Whittle-Matern with alpha = 1
##
Call:
graph_lme(formula = log_intensity ~ 1, graph = madrid_graph_new,
model = list(type = "WhittleMatern", alpha = 1), BC = 0)
##
Fixed effects:
Estimate Std.error z-value Pr(>|z|)
(Intercept) 6.1601 0.1421 43.36 <2e-16 ***
##
Random effects:
Estimate Std.error z-value
tau 0.41538 0.02282 18.205
kappa 0.63398 0.18992 3.338
##
```

```
Random effects (Matern parameterization):
Estimate Std.error z-value
sigma 2.1380 0.2856 7.487
range 3.1547 0.9201 3.429
##
Measurement error:
Estimate Std.error z-value
std. dev 0.71522 0.02888 24.77

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
Log-Likelihood: -1741.588
Number of function calls by 'optim' = 21
Optimization method used in 'optim' = L-BFGS-B
##
Time used to: fit the model = 31.84341 secs
```

## 7.1 Interpreting the Model Results

The estimated parameters of the Whittle-Matérn model provide insights into the spatial correlation of traffic intensity:

```
Extract parameters from the model
summary(fit_alpha1)
```

```
##
Latent model - Whittle-Matern with alpha = 1
##
Call:
graph_lme(formula = log_intensity ~ 1, graph = madrid_graph_new,
model = list(type = "WhittleMatern", alpha = 1), BC = 0)
##
Fixed effects:
Estimate Std.error z-value Pr(>|z|)
(Intercept) 6.1601 0.1421 43.36 <2e-16 ***
##
Random effects:
Estimate Std.error z-value
tau 0.41538 0.02282 18.205
kappa 0.63398 0.18992 3.338
##
Random effects (Matern parameterization):
Estimate Std.error z-value
sigma 2.1380 0.2856 7.487
range 3.1547 0.9201 3.429
##
Measurement error:
Estimate Std.error z-value
std. dev 0.71522 0.02888 24.77

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
Log-Likelihood: -1741.588
Number of function calls by 'optim' = 21
Optimization method used in 'optim' = L-BFGS-B
```

```
##
```

```
Time used to: fit the model = 31.84341 secs
```

- **Intercept:** The average log traffic intensity across the network
- **Range:** The effective distance (in kilometers) at which spatial correlation becomes negligible
- **Sigma:** The standard deviation of the spatial process, indicating the variability of traffic intensity
- **Nugget:** The measurement error or micro-scale variation

## 7.2 Predicting on the Graph


We can use our fitted model to make predictions at unobserved locations on the graph:

```
Build a mesh on the graph for prediction (if not already done)
if (length(madrid_graph_new$mesh$VtE) == 0) {
 madrid_graph_new$build_mesh(h = 0.01)
}

Create a data frame with the prediction locations
Extract mesh vertices as sf object
mesh_vertices_df <- data.frame(
 edge_number = madrid_graph_new$mesh$VtE[,1],
 distance_on_edge = madrid_graph_new$mesh$VtE[,2]
)

Make predictions
pred <- predict(fit_alpha1, newdata = mesh_vertices_df,
 normalize = TRUE)

Plot the predictions
madrid_graph_new$plot_function(X = pred$mean,
 vertex_size = 0,
 edge_width = 0.5) +
 scale_color_viridis_c(name = "Predicted\nLog Intensity")
```



## 7.3 Fitting the model with inlabru interface

Let us create the SPDE model object:

```
spde_model <- graph_spde(madrid_graph_new, alpha = 1)
```

Now, let us create the data object:

```
data_spde <- graph_data_spde(spde_model, loc_name = "loc")
```

Let us now fit the model:

```
f.s <- log_intensity ~ Intercept(1) +
 field(loc, model = spde_model)
fit_alpha1_inlabru <- bru(f.s,
 data = data_spde[["data"]])
```

Now, let us look at the model summary:

```
summary(fit_alpha1_inlabru)
```

```
inlabru version: 2.12.0
INLA version: 25.03.13
Components:
Intercept: main = linear(1), group = exchangeable(1L), replicate = iid(1L), NULL
field: main = cgeneric(loc), group = exchangeable(1L), replicate = iid(1L), NULL
Likelihoods:
Family: 'gaussian'
Tag: ''
Data class: 'metric_graph_data', 'list'
Response class: 'numeric'
Predictor: log_intensity ~ .
Used components: effects[Intercept, field], latent[]
Time used:
Pre = 1.9, Running = 1.3, Post = 0.173, Total = 3.38
Fixed effects:
mean sd 0.025quant 0.5quant 0.975quant mode kld
Intercept 6.008 0.162 5.663 6.014 6.316 6.014 0

Random effects:
Name Model
field CGeneric

Model hyperparameters:
mean sd 0.025quant 0.5quant 0.975quant mode
Precision for the Gaussian observations 1.98 0.160 1.685 1.976 2.315 1.963
Theta1 for field 0.88 0.054 0.772 0.881 0.984 0.884
Theta2 for field 1.86 0.472 1.033 1.830 2.876 1.684

Deviance Information Criterion (DIC): 3156.56
Deviance Information Criterion (DIC, saturated): 1747.61
Effective number of parameters: 529.61

Watanabe-Akaike information criterion (WAIC) ...: 3159.00
Effective number of parameters: 414.03

Marginal log-Likelihood: -1762.88
is computed
Posterior summaries for the linear predictor and the fitted values are computed
(Posterior marginals needs also 'control.compute=list(return.marginals.predictor=TRUE)')
```

Now, let us see the estimated parameters:

```
spde_result <- spde_metric_graph_result(fit_alpha1_inlabru,
 "field", spde_model)

summary(spde_result)
```

```
mean sd 0.025quant 0.5quant 0.975quant mode
sigma 3.11227 0.706029 2.05712 3.00056 4.8127 2.83778
range 7.20692 3.888480 2.82428 6.16900 17.5581 4.70193
```

## 8 Exercise 3: Further analysis

- Consider removing the zero values from the data.
- Try pruning the graph.
- Try fitting a model with smoothness parameter  $\alpha = 2$ .
- Try fitting the model estimating the smoothness parameter by using **rSPDE** interface.
- Try fitting the model with **INLA** interface.
- Check the connected components of the graph and consider whether more features should be added to the graph, or if the tolerances should be changed.